

# Text Technologies for Data Science: Assessment 3

s1107496

November 3, 2014

## 1 Introduction

This report describes the implementation of algorithms to detect duplicate (type 1), near duplicate (type 2), and number-heavy exact and near duplicate (type 3) news stories.

## 2 Type 1 duplicate detection (task 2)

Through manual examination, type 1 duplicates were determined to be exact duplicates. To detect these, a hash table with keys consisting of untokenized stories (with IDs removed) and values consisting of sets of story IDs was maintained. If, for a given key, there existed a matching story, the corresponding IDs were written to file. Otherwise, the story/ID key/value pair for the current line was written to the hash table.

## 3 Type 2 duplicate detection (task 3)

### 3.1 Implementation overview

To detect type 2 duplicates, stories were first tokenized by splitting on whitespace and other non-alphanumeric characters, using the regular expression suggested by Victor Lavrenko. Stories that had already been detected as type 1 duplicates were skipped. Stopwords, as given by `nltk`, were filtered out before frequencies for each token were computed. A simhash implementation using the md5 hashing algorithm from `hashlib` was then used to find potential duplicates, with tokens weighted by frequency. A difference metric between these potential duplicates was then computed by subtracting the size of the union of their token sets from the size of the intersection of their token sets; the first (if any) duplicate with fewer than three differences was written to file.

### 3.2 K and L optimization

While the initial values of 33 and 3 chosen for K and L (respectively) achieved 100% recall on `data.train`, these values resulted in a running time on `data.test` of over 30 minutes. To remedy this, values for K and L were hand-optimized to maximize expected recall on `data.test` (see figure below). With K set to 28 and L set to 1, 66.67% of items from `data.train` were correctly retrieved, while also processing XXX% of `data.test` within 30 minutes (as measured on a DICE machine with an i5 processor and 8GB RAM). Because this was not the entirety of `data.test`, the decision was made to stop processing lines from `data.test` after 29.5 minutes.

## 4 Comparison of algorithm running times

Tests were run on a DICE machine with an i5 processor and 8GB RAM.